Graphcore

Initial results (running mnist_poptorch.py as is):

```
TrainingModelWithLoss(
  (model): Network(
    (layer1): Block(
      (conv): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1))
      (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
      (relu): ReLU()
    )
    (layer2): Block(
      (conv): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
      (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
      (relu): ReLU()
    )
    (layer3): Linear(in_features=1600, out_features=128, bias=True)
    (layer3_act): ReLU()
    (layer3_dropout): Dropout(p=0.5, inplace=False)
    (layer4): Linear(in_features=128, out_features=10, bias=True)
    (softmax): Softmax(dim=1)
  )
  (loss): CrossEntropyLoss()
)
Accuracy on test set: 98.62%
```

Changing batch size:

Batch_size = 24 (below)

```
    (model): Network(
      (layer1): Block(
        (conv): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1))
        (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=Fals
        (relu): ReLU()
      )
      (layer2): Block(
        (conv): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=Fals
        (relu): ReLU()
      )
      (layer3): Linear(in_features=1600, out_features=128, bias=True)
      (layer3_act): ReLU()
      (layer3_dropout): Dropout(p=0.5, inplace=False)
      (layer4): Linear(in_features=128, out_features=10, bias=True)
      (softmax): Softmax(dim=1)
    )
    (loss): CrossEntropyLoss()
  )
  Accuracy on test set: 98.05%
  (poptorch33_env) khalil_nor@gc-poplar-03:~/graphcore/examples/tutorials/simple_applica
-login-01.ai.alcf.anl.gov  ⊗ 0 ⚠ 0  0
```

Batch_size = 56 (below)

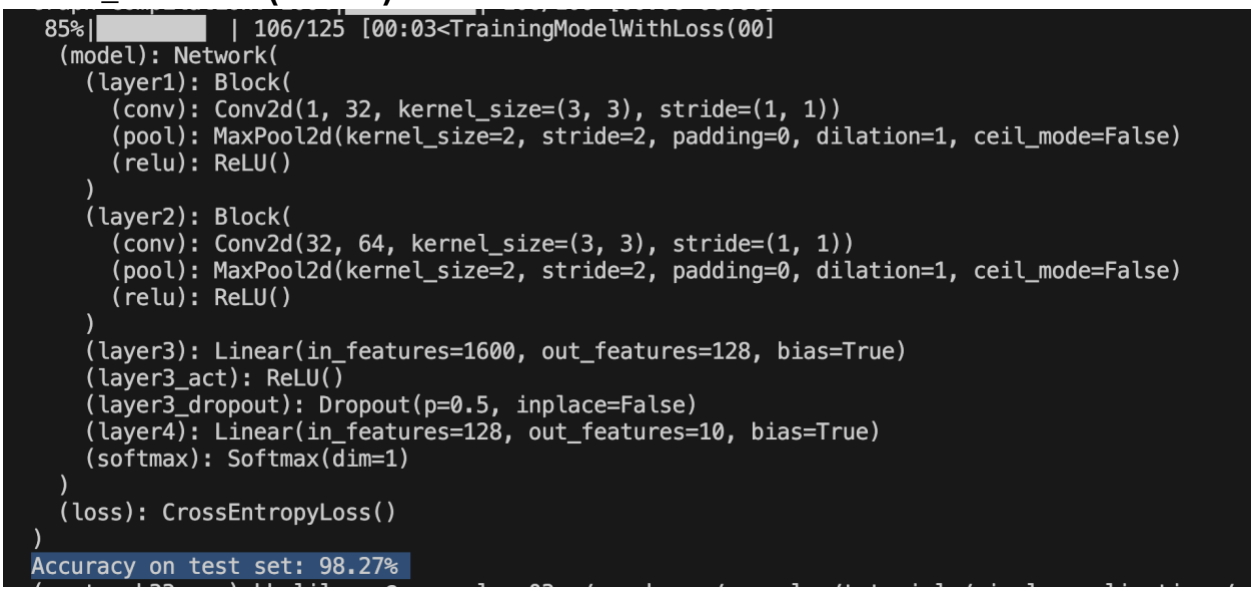

Batch_size = 128

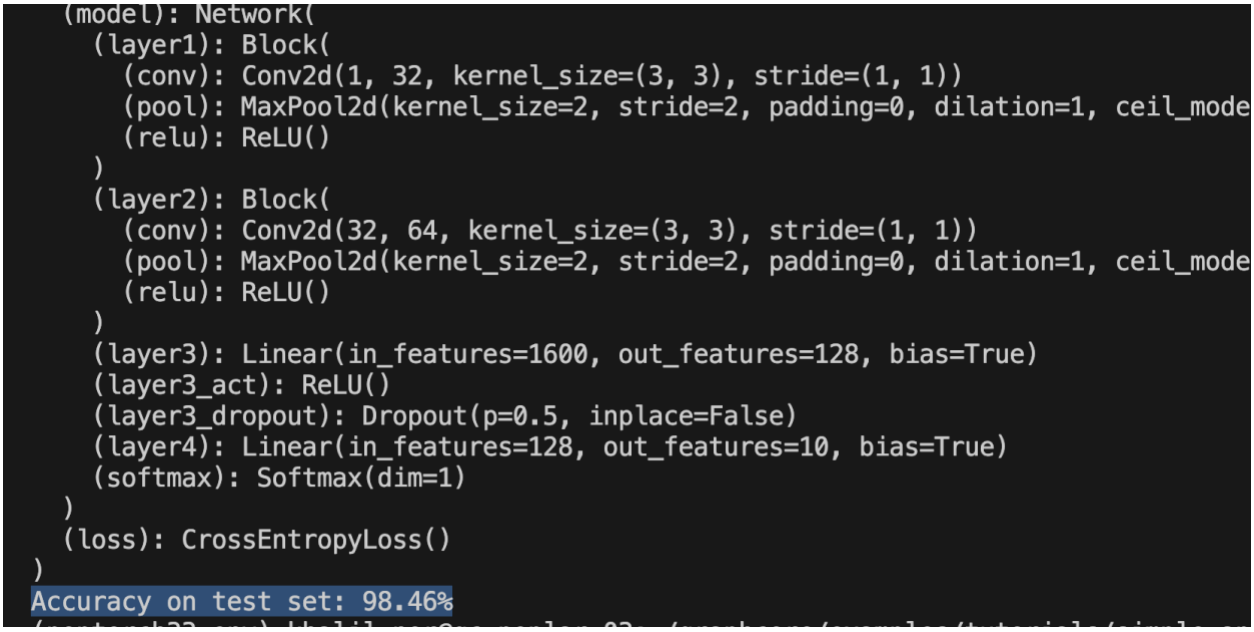

When set to its initial batch size of 8, the model exhibited the highest test accuracy (98.62%). Generally, except for the initial case, increasing the batch size from 24 to 56 to 128 lead to small increases in test accuracy, although these were very minimal.

Changing learning rate:

Lr = 0.06

```
Graph compilation: 100%|          | 100/100 [00:00<00:00]
Epochs: 100%|          | 10/10 [01:26<00:00,  8.67s/it]
Graph compilation: 100%|          | 100/100 [00:00<00:00]
 86%|          | 107/125 [00:TrainingModelWithLoss(00:00]
  (model): Network(
    (layer1): Block(
      (conv): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1))
      (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=F
      (relu): ReLU()
    )
    (layer2): Block(
      (conv): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
      (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=F
      (relu): ReLU()
    )
    (layer3): Linear(in_features=1600, out_features=128, bias=True)
    (layer3_act): ReLU()
    (layer3_dropout): Dropout(p=0.5, inplace=False)
    (layer4): Linear(in_features=128, out_features=10, bias=True)
    (softmax): Softmax(dim=1)
  )
  (loss): CrossEntropyLoss()
)
Accuracy on test set: 97.84%
```

Lr = 0.12

```
  (model): Network(
    (layer1): Block(
      (conv): Conv2d(1, 32, kernel_size=(3, 3), stride=(1, 1))
      (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
      (relu): ReLU()
    )
    (layer2): Block(
      (conv): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
      (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
      (relu): ReLU()
    )
    (layer3): Linear(in_features=1600, out_features=128, bias=True)
    (layer3_act): ReLU()
    (layer3_dropout): Dropout(p=0.5, inplace=False)
    (layer4): Linear(in_features=128, out_features=10, bias=True)
    (softmax): Softmax(dim=1)
  )
  (loss): CrossEntropyLoss()
)
Accuracy on test set: 98.15%
```

Lr = 0.25

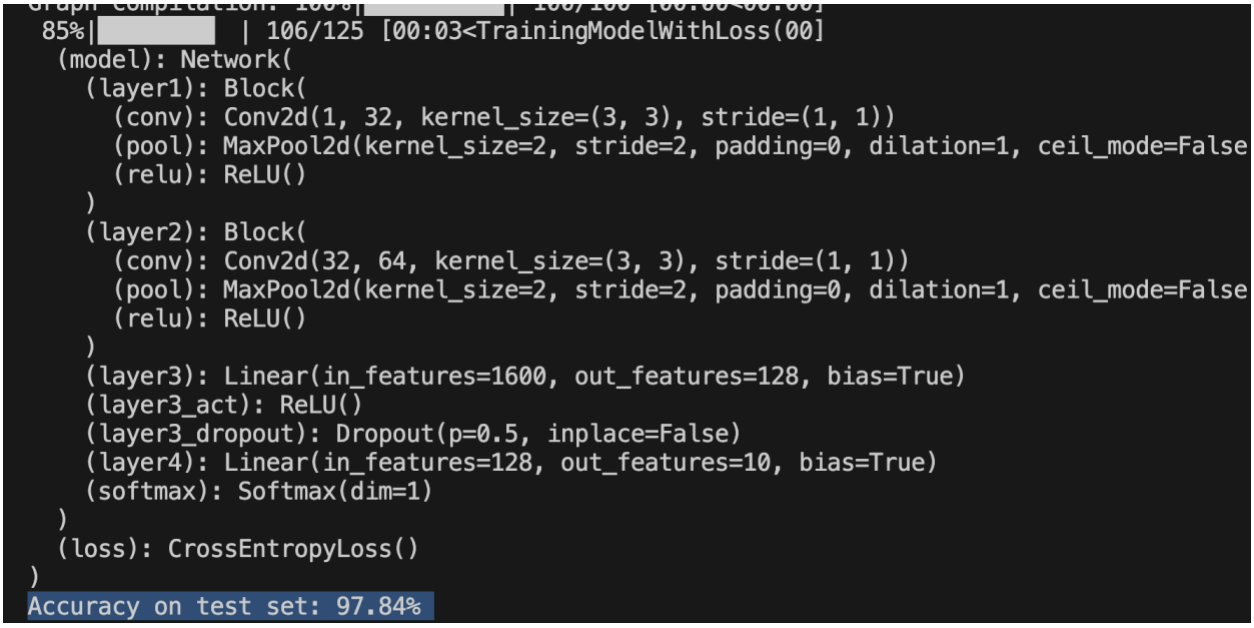

Lr = 0.35

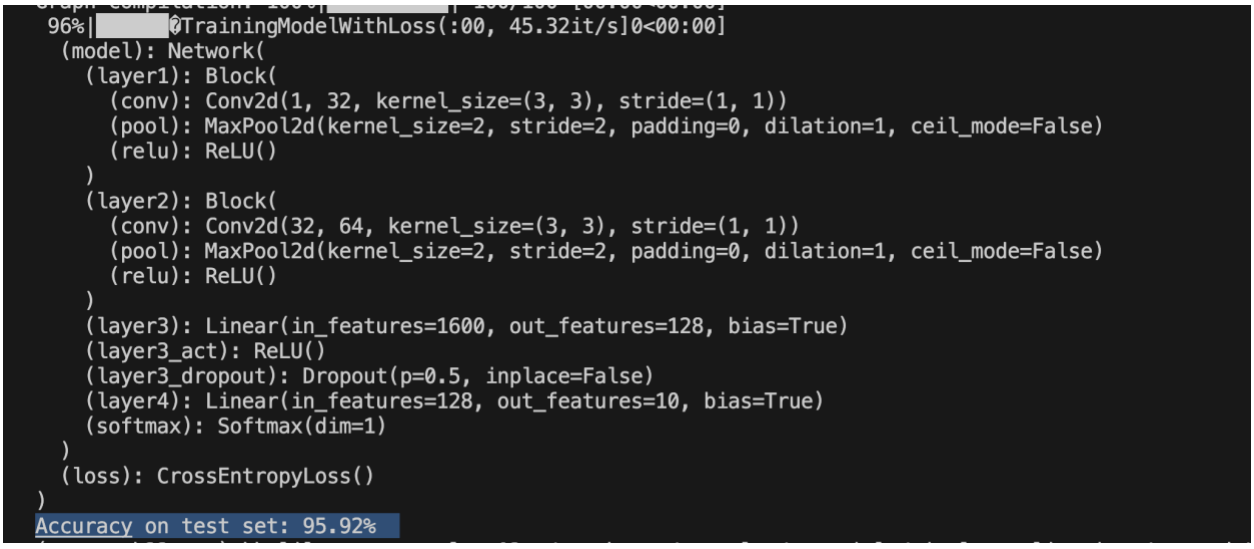

Greatly increasing the learning rate significantly decreases the accuracy on the test set.

Changing number of epochs:
Epochs = 4

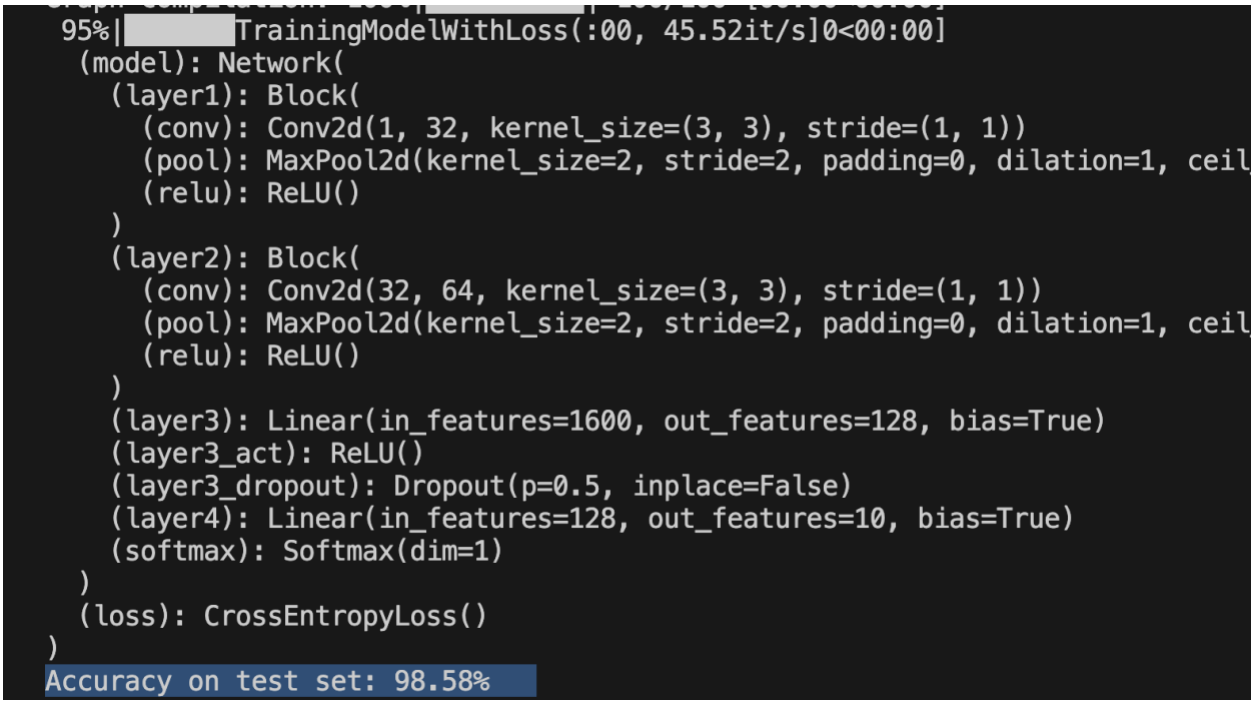

Epochs = 7

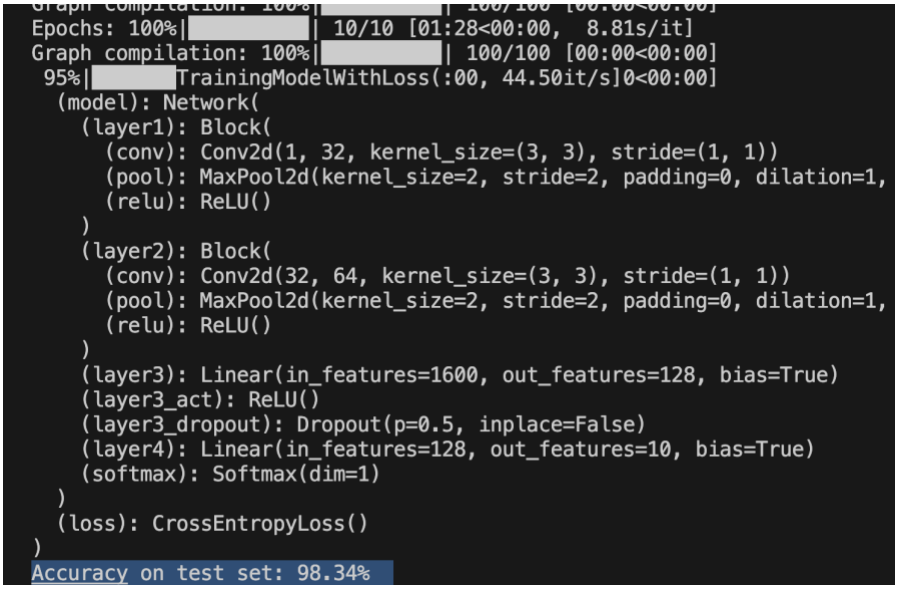

Epochs = 12

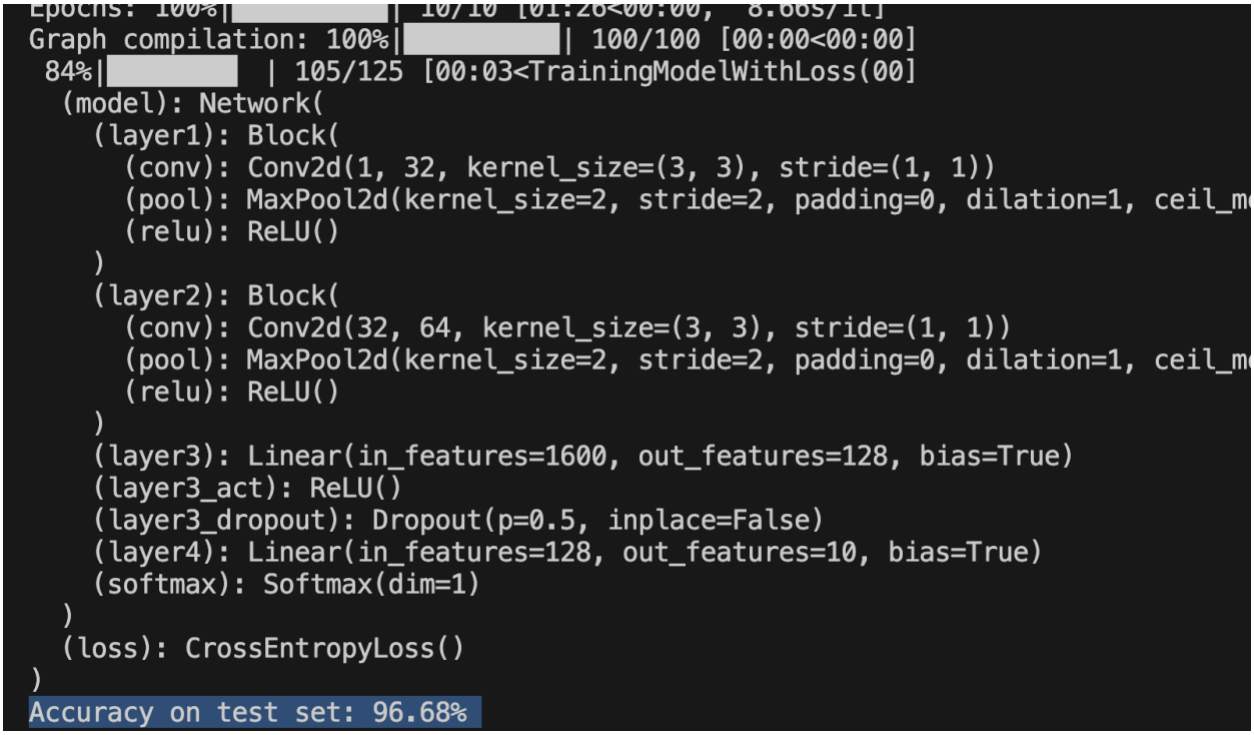

Epochs = 20

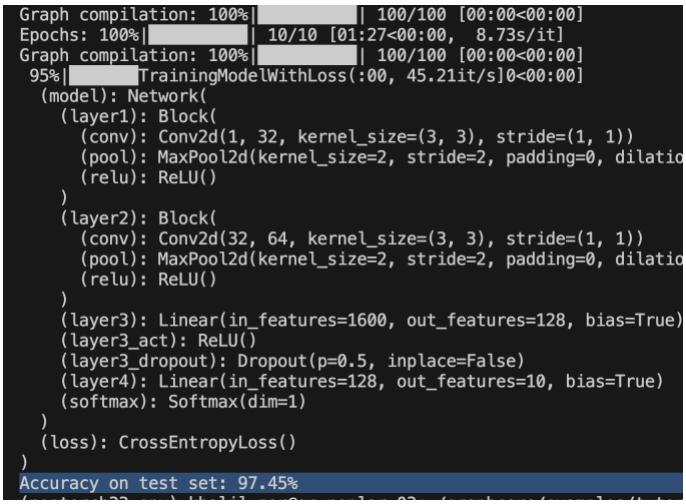

Increasing to be > 10 resulted in a decreased test accuracy, while keeping the epochs between 4 – 10 resulted in a test accuracy that was basically the same.